

# Scrape the Web

Strategies for programming websites that don't expect it

Asheesh Laroia [asheesh@asheesh.org](mailto:asheesh@asheesh.org)

January 18, 2012

# This is supposed to be practice

- ▶ But there are things I know I'll have to fix

# This is supposed to be practice

- ▶ But there are things I know I'll have to fix
- ▶ More graphs and diagrams

# This is supposed to be practice

- ▶ But there are things I know I'll have to fix
- ▶ More graphs and diagrams
  - ▶ But which ones?

# This is supposed to be practice

- ▶ But there are things I know I'll have to fix
- ▶ More graphs and diagrams
  - ▶ But which ones?
- ▶ Better presentation style (sorry!)

So please do interrupt me

...and we begin

# Welcome



# Meta

- ▶ You will learn neat tricks

# Meta

- ▶ You will learn neat tricks
- ▶ DO NOT BECOME AN EVIL COMMENT SPAMMER

# Meta

- ▶ You will learn neat tricks
- ▶ DO NOT BECOME AN EVIL COMMENT SPAMMER
- ▶ Theory and running code

# Meta

- ▶ You will learn neat tricks
- ▶ DO NOT BECOME AN EVIL COMMENT SPAMMER
- ▶ Theory and running code
- ▶ Brittle? Sometimes.

# Meta

- ▶ You will learn neat tricks
- ▶ DO NOT BECOME AN EVIL COMMENT SPAMMER
- ▶ Theory and running code
- ▶ Brittle? Sometimes.
- ▶ The comics aren't mine; ask for references.

# Things you'll need

- ▶ Sample code: <http://FIXME.com>
- ▶ Install FireBug
- ▶ Install python-lxml if it is easy

# Pacing

- ▶ Slow me down,

# Pacing

- ▶ Slow me down,
- ▶ or speed me up.
- ▶ With your voice, or by raising your hand.



# What is web scraping?

Generally speaking,

- ▶ You retrieve some data from the web,

# What is web scraping?

Generally speaking,

- ▶ You retrieve some data from the web,
- ▶ You extract some information,

# What is web scraping?

Generally speaking,

- ▶ You retrieve some data from the web,
- ▶ You extract some information,
- ▶ and optionally you repeat.

# Perspectives on scraping

- ▶ One page vs. a whole site

# Perspectives on scraping

- ▶ One page vs. a whole site
- ▶ A site's contents now, or for the future as well

## (1) Diving in with curry



# Lunchtime

- ▶ <http://mehfilindian.com/LunchMenuTakeOut.htm>



# Lunchtime

- ▶ <http://mehfilindian.com/LunchMenuTakeOut.htm>
- ▶ A question

# Lunchtime

- ▶ <http://mehfilindian.com/LunchMenuTakeOut.htm>
- ▶ A question
  - ▶ is there eggplant today?

# From python

examples/curry/trivial.py

## (2) HTML: Structured text on the web

# Two easy ways to read HTML

In a browser:

- ▶ View source

# Two easy ways to read HTML

In a browser:

- ▶ View source
- ▶ Inspect element (requires Firebug or DOM Inspector)

# HTML vs. XHTML

- ▶ Both are trees of tags

# HTML vs. XHTML

- ▶ Both are trees of tags
- ▶ HTML: from 1992



# HTML vs. XHTML

- ▶ Both are trees of tags
- ▶ HTML: from 1992
- ▶ XHTML: from 2000

# HTML vs. XHTML

- ▶ Both are trees of tags
- ▶ HTML: from 1992
- ▶ XHTML: from 2000
- ▶ ...did XHTML win?

# Stats pop quiz: size and type

(Stats from the MAMA survey published by Opera,  
<http://dev.opera.com/articles/view/mama-key-findings/>.)

- ▶ Average page size?

# Stats pop quiz: size and type

(Stats from the MAMA survey published by Opera,  
<http://dev.opera.com/articles/view/mama-key-findings/>.)

- ▶ Average page size?
  - ▶ 16.5K

# Stats pop quiz: size and type

(Stats from the MAMA survey published by Opera,  
<http://dev.opera.com/articles/view/mama-key-findings/>.)

- ▶ Average page size?
  - ▶ 16.5K
- ▶ HTML to XHTML ratio?

# Stats pop quiz: size and type

(Stats from the MAMA survey published by Opera,  
<http://dev.opera.com/articles/view/mama-key-findings/>.)

- ▶ Average page size?
  - ▶ 16.5K
- ▶ HTML to XHTML ratio?
  - ▶ 2:1

# Stats pop quiz: quirks and tags (pt. 1)

- ▶ Transitional vs. Strict+Framset?

# Stats pop quiz: quirks and tags (pt. 1)

- ▶ Transitional vs. Strict+Framset?
  - ▶ 10:1



## Stats pop quiz: quirks and tags (pt. 1)

- ▶ Transitional vs. Strict+Frameset?
  - ▶ 10:1
- ▶ How many pages render in "Quirks" mode?

# Stats pop quiz: quirks and tags (pt. 1)

- ▶ Transitional vs. Strict+Frameset?
  - ▶ 10:1
- ▶ How many pages render in "Quirks" mode?
  - ▶ 85 percent

# Stats pop quiz: quirks and tags (pt. 1)

- ▶ Transitional vs. Strict+Frameset?
  - ▶ 10:1
- ▶ How many pages render in "Quirks" mode?
  - ▶ 85 percent
- ▶ What percent validate?

# Stats pop quiz: quirks and tags (pt. 1)

- ▶ Transitional vs. Strict+Frameset?
  - ▶ 10:1
- ▶ How many pages render in "Quirks" mode?
  - ▶ 85 percent
- ▶ What percent validate?
  - ▶ 4.13 percent

## Stats pop quiz: quirks and tags (pt. 2)

- ▶ What's more popular? TITLE or BODY?
- ▶ The web is a mess

## Stats pop quiz: quirks and tags (pt. 2)

- ▶ What's more popular? TITLE or BODY?
  - ▶ TITLE
  
- ▶ The web is a mess

## Stats pop quiz: quirks and tags (pt. 2)

- ▶ What's more popular? TITLE or BODY?
  - ▶ TITLE
- ▶ What percent of web pages with validations badges actually validate?
- ▶ The web is a mess

## Stats pop quiz: quirks and tags (pt. 2)

- ▶ What's more popular? TITLE or BODY?
  - ▶ TITLE
- ▶ What percent of web pages with validations badges actually validate?
  - ▶ about 50 percent
- ▶ The web is a mess



### (3) Parsing "HTML" in Python

# Things to consider

- ▶ What does a parser do with invalid HTML?

# Things to consider

- ▶ What does a parser do with invalid HTML?
- ▶ Does it handle XHTML properly?

# Things to consider

- ▶ What does a parser do with invalid HTML?
- ▶ Does it handle XHTML properly?
  - ▶ They all do; don't worry.

# Things to consider

- ▶ What does a parser do with invalid HTML?
- ▶ Does it handle XHTML properly?
  - ▶ They all do; don't worry.
- ▶ `examples/parsing/` has samples.

# A showcase of some options

- ▶ HTMLParser (stdlib!)
- ▶ xml.dom.minidom (stdlib!)

# A showcase of some options

- ▶ HTMLParser (stdlib!)
- ▶ xml.dom.minidom (stdlib!)
- ▶ BeautifulSoup
- ▶ html5lib
- ▶ lxml.html

## (4) On regular expressions



- ▶ Some people, when confronted with a problem, think
  - ▶ "I know, I'll use regular expressions."
- ▶ Now they have two problems. – jwz.

# But why?

- ▶ `a href="whatever"`
- ▶ `a href='whatever'`
- ▶ `a href="whatever'`

But it's good enough for...

- ▶ Curry

# But it's good enough for...

- ▶ Curry
- ▶ Text analysis:

# But it's good enough for...

- ▶ Curry
- ▶ Text analysis:
  - ▶ Reviews 1-10 of 430

## If you have to

- ▶ Use Kodos, a regular expression GUI
- ▶ (Note: redemo.py in Python source is unmaintained.)

# If you have to

- ▶ Use Kodos, a regular expression GUI
- ▶ (Note: redemo.py in Python source is unmaintained.)
- ▶ Be conservative in what you do, be liberal in what you accept from others. – Jon Postel.

## (5) Parsers in depth



# Searching document trees

- ▶ BeautifulSoup API  
([examples/tree-builders/beautifulsoup/search.py](https://examples/tree-builders/beautifulsoup/search.py))

# Searching document trees

- ▶ BeautifulSoup API  
([examples/tree-builders/beautifulsoup/search.py](https://www.crummy.com/software/BeautifulSoup/bs4/doc/#searching))
- ▶ html5lib creates BeautifulSoup objects (or others)  
([examples/tree-builders/html5lib/search.py](https://www.crummy.com/software/BeautifulSoup/bs4/doc/#searching))

# Searching document trees

- ▶ BeautifulSoup API  
([examples/tree-builders/beautifulsoup/search.py](#))
- ▶ html5lib creates BeautifulSoup objects (or others)  
([examples/tree-builders/html5lib/search.py](#))
- ▶ lxml provides XPath ([examples/tree-builders/lxml/searchxpath.py](#))

# Searching document trees

- ▶ BeautifulSoup API  
([examples/tree-builders/beautifulsoup/search.py](#))
- ▶ html5lib creates BeautifulSoup objects (or others)  
([examples/tree-builders/html5lib/search.py](#))
- ▶ lxml provides XPath ([examples/tree-builders/lxml/searchxpath.py](#))
- ▶ lxml provides CSSSelect ([examples/tree-builders/lxml/searchcss.py](#))

# General structure of scraping a page

- ▶ Get a page's HTML

# General structure of scraping a page

- ▶ Get a page's HTML
- ▶ Parse it

# General structure of scraping a page

- ▶ Get a page's HTML
- ▶ Parse it
- ▶ Pull out the items you need

# General structure of scraping a page

- ▶ Get a page's HTML
- ▶ Parse it
- ▶ Pull out the items you need
- ▶ Return them as a dictionary, or an object



# A closer look at curry

- ▶ (see Python interpreter)
- ▶ (let's use BeautifulSoup)

# A closer look at curry

- ▶ (see Python interpreter)
- ▶ (let's use BeautifulSoup)
- ▶ Conclusion: This is a text-processing problem, not a tag problem.

# Mini-lesson

Three kinds of page:

- ▶ Hand-written pages

# Mini-lesson

Three kinds of page:

- ▶ Hand-written pages
- ▶ Machine-written pages

# Mini-lesson

Three kinds of page:

- ▶ Hand-written pages
- ▶ Machine-written pages
- ▶ Machine-written pages, old-skool

# More BeautifulSoup: Yahoo! Finance

- ▶ `examples/tree-builders/beautifulsoup_yfinance.py`

## (6) Interacting with the web

# Hard-coding URLs: Yahoo! search

- ▶ `examples/search/yahoo.py`
- ▶ `examples/search/google.py`



# More about HTTP: Headers

- ▶ (Firefox demo)

# More about HTTP: Status codes

- ▶ 2xx: Success
- ▶ 3xx: Redirection
- ▶ 4xx: Error

# More about HTTP: Status codes

- ▶ 2xx: Success
- ▶ 3xx: Redirection
- ▶ 4xx: Error
- ▶ 402: Payment Required
- ▶ 404 Not Found

# More about HTTP: Status codes

- ▶ 2xx: Success
- ▶ 3xx: Redirection
- ▶ 4xx: Error
- ▶ 402: Payment Required
- ▶ 404 Not Found
- ▶ 410 Gone

# More about HTTP: Status codes

- ▶ 2xx: Success
- ▶ 3xx: Redirection
- ▶ 4xx: Error
- ▶ 402: Payment Required
- ▶ 404 Not Found
- ▶ 410 Gone
- ▶ 418 I'm a teapot

# More about HTTP: Methods

- ▶ GET
- ▶ POST
- ▶ PUT

# More about HTTP: Methods

- ▶ GET
- ▶ POST
- ▶ PUT
- ▶ BREW

# Once we set User-Agent, are we just like Firefox?

- ▶ JavaScript behavior



# Once we set User-Agent, are we just like Firefox?

- ▶ JavaScript behavior
- ▶ Image download behavior

# Once we set User-Agent, are we just like Firefox?

- ▶ JavaScript behavior
- ▶ Image download behavior
- ▶ Cookie behavior

# Once we set User-Agent, are we just like Firefox?

- ▶ JavaScript behavior
- ▶ Image download behavior
- ▶ Cookie behavior
- ▶ Invalid HTML handling behavior (?)

# Once we set User-Agent, are we just like Firefox?

- ▶ JavaScript behavior
- ▶ Image download behavior
- ▶ Cookie behavior
- ▶ Invalid HTML handling behavior (?)
- ▶ Accept: headers

# Google, again

- ▶ `examples/search/urllib2-user-agent/google as ie.py`

# Google, again

- ▶ `examples/search/urllib2-user-agent/google as ie.py`
- ▶ IE 5 vs. IE 8

# robots.txt

- ▶ User-agent: \*
- ▶ Disallow: /
- ▶ Allow: /crawlme.html

# robots.txt

- ▶ User-agent: \*
- ▶ Disallow: /
- ▶ Allow: /crawlme.html
- ▶ <http://www.robotstxt.org/>



# robots.txt

- ▶ User-agent: \*
- ▶ Disallow: /
- ▶ Allow: /crawlme.html
- ▶ <http://www.robotstxt.org/>
- ▶ Don't ever GET it

(7) Filling out forms, and handling cookies, with mechanize

# The weather (by hand)

- ▶ <http://cepstral.com/cgi-bin/demos/weather>

# The weather (by hand)

- ▶ <http://cepstral.com/cgi-bin/demos/weather>
- ▶ Find the POST target in Firebug

# The weather (by hand)

- ▶ <http://cepstral.com/cgi-bin/demos/weather>
- ▶ Find the POST target in Firebug
- ▶ [examples/cepstral/just post.py](#)

# The weather (by hand)

- ▶ <http://cepstral.com/cgi-bin/demos/weather>
- ▶ Find the POST target in Firebug
- ▶ [examples/cepstral/just post.py](#)
- ▶ [examples/cepstral/play wav.py](#)

# The weather (with mechanize)

- ▶ `examples/cepstral/just post via mechanize.py`

# Search the web (via mechanize)

- ▶ `examples/search/yahoo mechanize.py`



# Search the web (via mechanize)

- ▶ `examples/search/yahoo mechanize.py`
- ▶ Oh snap, we're a robot.

# Search the web (via mechanize)

- ▶ `examples/search/yahoo mechanize.py`
- ▶ Oh snap, we're a robot.
- ▶ `examples/search/yahoo mechanize norobots.py`

# Search the web (via mechanize)

- ▶ `examples/search/yahoo mechanize.py`
- ▶ Oh snap, we're a robot.
- ▶ `examples/search/yahoo mechanize norobots.py`
- ▶ `examples/search/google mechanize.py`

## (8) Recap and philosophy

# Things we've seen

- ▶ Loading web pages from the network with urllib2

# Things we've seen

- ▶ Loading web pages from the network with urllib2
- ▶ Parsing web pages (even broken ones)

# Things we've seen

- ▶ Loading web pages from the network with urllib2
- ▶ Parsing web pages (even broken ones)
- ▶ Scraping that page into a set of structured Python objects

# Things we've seen

- ▶ Loading web pages from the network with urllib2
- ▶ Parsing web pages (even broken ones)
- ▶ Scraping that page into a set of structured Python objects
- ▶ HTTP status codes



# Things we've seen

- ▶ Loading web pages from the network with urllib2
- ▶ Parsing web pages (even broken ones)
- ▶ Scraping that page into a set of structured Python objects
- ▶ HTTP status codes
- ▶ Faking the user agent header

# Things we've seen

- ▶ Loading web pages from the network with urllib2
- ▶ Parsing web pages (even broken ones)
- ▶ Scraping that page into a set of structured Python objects
- ▶ HTTP status codes
- ▶ Faking the user agent header
- ▶ Submitting forms

# Things we've seen

- ▶ Loading web pages from the network with urllib2
- ▶ Parsing web pages (even broken ones)
- ▶ Scraping that page into a set of structured Python objects
- ▶ HTTP status codes
- ▶ Faking the user agent header
- ▶ Submitting forms
- ▶ Keeping a session with cookies

## (9) Even more about parsers

# Things to look for

- ▶ Performance
- ▶ Ease-of-use
- ▶ Quality
- ▶ Maintained-ness

# Redux

- ▶ HTMLParser (stdlib!)

# Redux

- ▶ `xml.dom.minidom` (stdlib!)

# Redux

- ▶ `xml.dom.minidom` (stdlib!)
- ▶ BeautifulSoup



# Redux

- ▶ `xml.dom.minidom` (stdlib!)
- ▶ BeautifulSoup
- ▶ `html5lib`

# Redux

- ▶ `xml.dom.minidom` (stdlib!)
- ▶ BeautifulSoup
- ▶ `html5lib`
- ▶ `lxml.html`

# Winners

- ▶ Resilience: lxml.html == html5lib & BeautifulSoup & stdlib

# Winners

- ▶ Resilience: lxml.html == html5lib > BeautifulSoup > stdlib
- ▶ Performance: lxml.html > stdlib > BeautifulSoup > html5lib

## (10) Countermeasures

# The basics

- ▶ Referer header
- ▶ Cookies
- ▶ Hidden form fields

# The basics

- ▶ Referer header
- ▶ Cookies
- ▶ Hidden form fields
- ▶ Solved by mechanize

# The hard ones

- ▶ Per-IP address query limits



# The hard ones

- ▶ Per-IP address query limits
- ▶ Behavior profiling

# The hard ones

- ▶ Per-IP address query limits
- ▶ Behavior profiling
- ▶ JavaScript

# The hard ones

- ▶ Per-IP address query limits
- ▶ Behavior profiling
- ▶ JavaScript
- ▶ CAPTCHAs

# IP addresss

- ▶ ssh -D
- ▶ tsocks
- ▶ socks monkey

# IP addresss

- ▶ ssh -D
- ▶ tsocks
- ▶ socks monkey
- ▶ (All in the sample code)

# Behavior profiling

- ▶ You're doomed.

# JavaScript

Options:

1. Re-write the JS in Python
2. Send the JS to python-spidermonkey

# CAPTCHAs

- ▶ Hope for an easy one
  - ▶ <http://www.mailinator.com/images/captcha1.gif>



# CAPTCHAs

- ▶ Hope for an easy one
  - ▶ <http://www.mailinator.com/images/captcha1.gif>
- ▶ Otherwise, just present it to the human operator...

(11) When (and how) to just automate Firefox

# Selenium and friends

- ▶ `examples/seleniumrc/test google.py`
- ▶ Selenium RC for XPath

# Thanks

asheesh@asheesh.org